

*annotated
version*

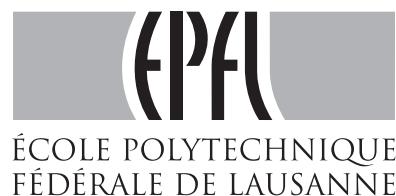


Machine Learning Course - CS-433

Optimization

Sep 27+29, 2016

©Martin Jaggi and Mohammad Emtiyaz Khan 2016



Learning / Estimation/ Fitting

Given a cost function $\mathcal{L}(\mathbf{w})$, we wish to find \mathbf{w}^* that minimizes the cost:

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) \quad \text{subject to } \mathbf{w} \in \mathbb{R}^D$$

This means the *learning* problem is formulated as an *optimization problem*.

We will use an *optimization algorithm* to solve the problem (to find a good \mathbf{w}).

Grid Search

Grid search is one of the simplest optimization algorithms. We compute the cost over all values \mathbf{w} in a grid, and pick the best among those.

This is brute-force, but extremely simple and works for any kind of cost function when we have very few parameters and the cost is easy to compute.

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (y_n - f_w(x_n))^2$$

model

Examples:

- linear Regression

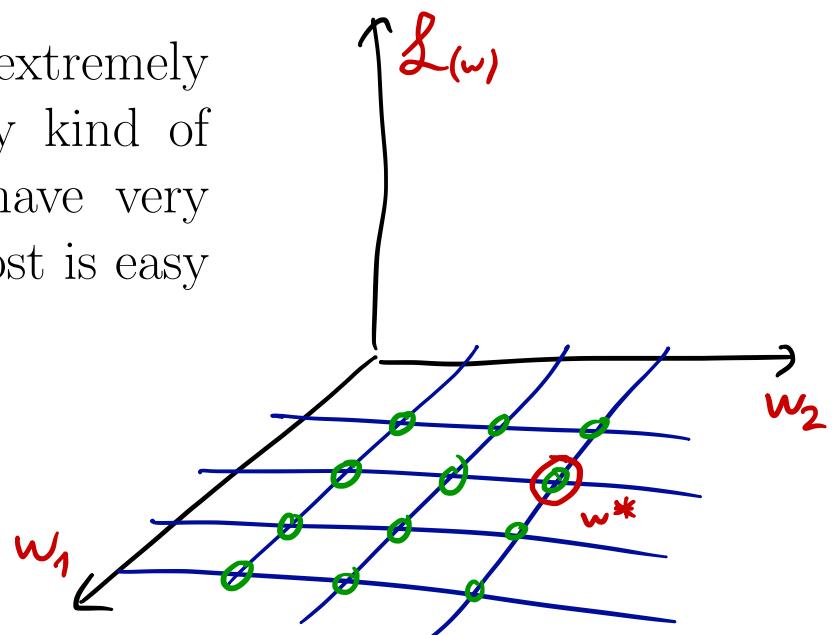
$$f_w(x) = x^T w$$

—.1

- Neural Network

$$f_w(x) = \text{[Diagram of a neural network with inputs } x_1, x_2, \dots, x_d \text{ and hidden layers with nodes and weights.]}$$

↑



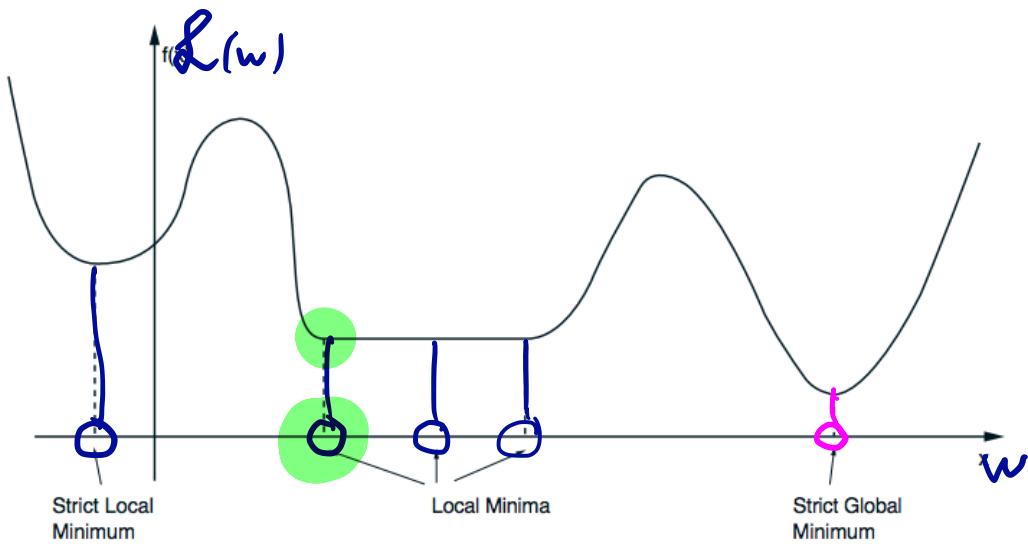
For a large number of parameters D , however, grid search has too many “for-loops”, resulting in an exponential computational complexity:

If we decide to use 10 possible values for each dimension of \mathbf{w} , then we have to check 10^D points. This is clearly impossible for most practical machine learning models, which can often have $D \approx$ millions of parameters. Choosing a good range of values for each dimension is another problem.

Are there any other issues?

Optimization Landscapes

find $w \in \mathbb{R}^D$
 $\min_w \mathcal{L}(w)$



The above figure is taken from Bertsekas, Nonlinear programming.

A vector w^* is a local minimum of \mathcal{L} if it is no worse than its neighbors; i.e. there exists an $\epsilon > 0$ such that,

$$\mathcal{L}(w^*) \leq \mathcal{L}(w), \quad \forall w \text{ with } \|w - w^*\| < \epsilon$$

A vector w^* is a global minimum of \mathcal{L} if it is no worse than all others,

$$\mathcal{L}(w^*) \leq \mathcal{L}(w), \quad \forall w \in \mathbb{R}^D$$

These local or global minimum are said to be strict if the corresponding inequality is strict for $w \neq w^*$.

When working with a cost function, it is essential to make sure that a global minimum exists, i.e. that the cost function is lower bounded.

=differentiable

Smooth Optimization

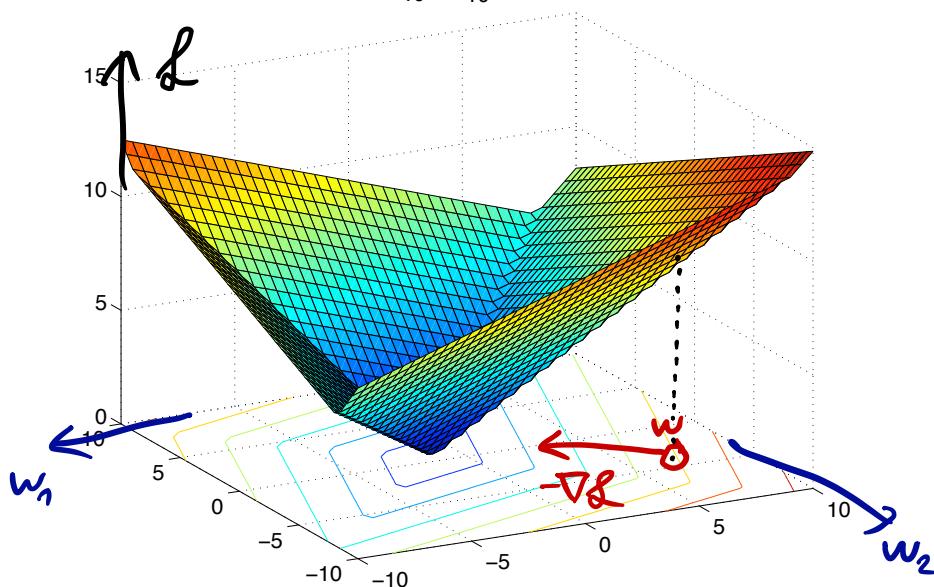
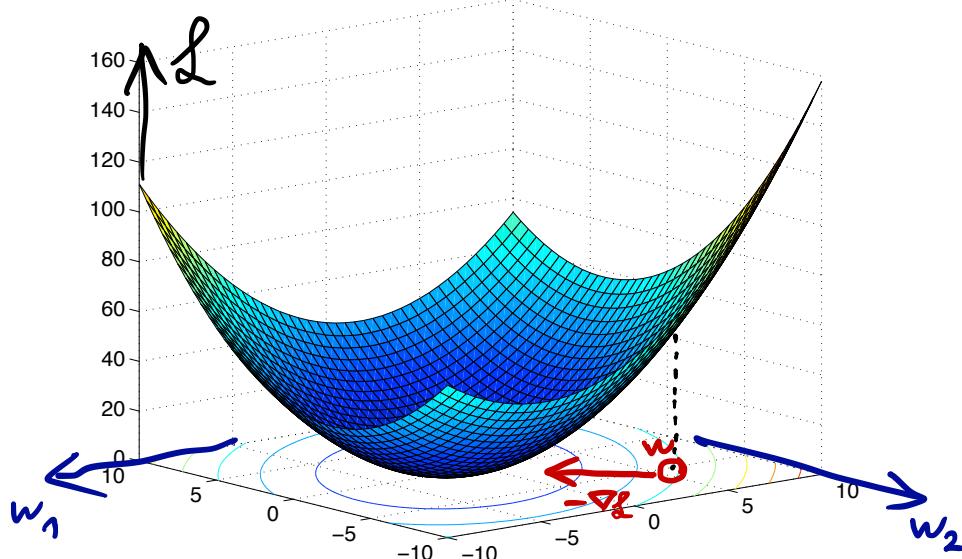
Follow the Gradient

A gradient (at a point) is the slope of the tangent to the function (at that point). It points to the direction of largest increase of the function.

For a 2-parameter model, $\text{MSE}(\mathbf{w})$

and $\text{MAE}(\mathbf{w})$ are shown below.

(We used $\mathbf{y}_n \approx w_0 + w_1 x_{n1}$ with $\mathbf{y}^\top = [2, -1, 1.5]$ and $\mathbf{x}^\top = [-1, 1, -1]$).



Definition of the gradient:

$$\nabla \mathcal{L}(\mathbf{w}) := \left[\frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_D} \right]^\top$$

This is a vector, $\nabla \mathcal{L}(\mathbf{w}) \in \mathbb{R}^D$.

how quickly does \mathcal{L} change depending on variable w_1 ?

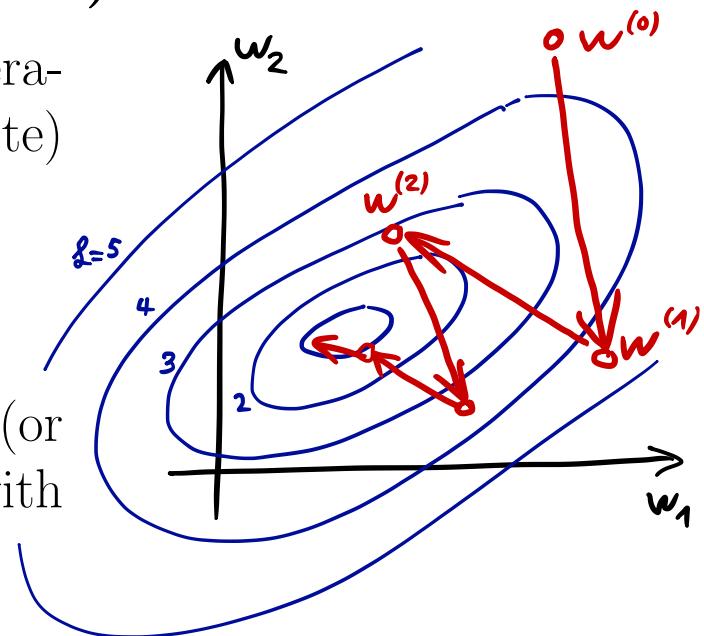
same for w_D

Gradient Descent (Batch)

To minimize the function, we iteratively take a step in the (opposite) direction of the gradient

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

where $\gamma > 0$ is the step-size (or learning rate). Then repeat with the next t .



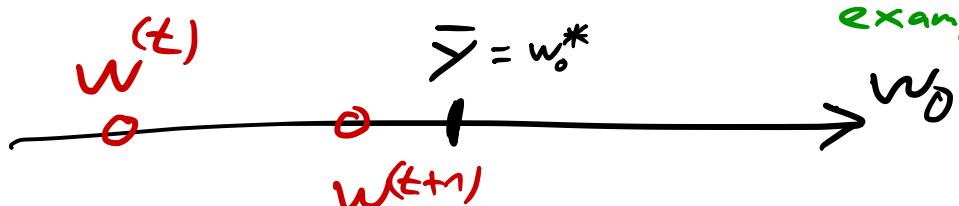
Example: Gradient descent for 1-parameter model to minimize MSE:

$$w_0^{(t+1)} := (1 - \gamma)w_0^{(t)} + \gamma \bar{y}$$

Where $\bar{y} = \sum_n y_n / N$. When is this sequence guaranteed to converge?

$$\gamma \in (0, 1]$$

→ convergence in this example



$\gamma > 1$ too?

Gradient Descent for Linear MSE

For linear regression

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}, \mathbf{X} = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1D} \\ x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \dots & x_{ND} \end{bmatrix}$$

We define the error vector \mathbf{e} :

$$\mathbf{e} := \mathbf{y} - \mathbf{X}\mathbf{w}$$

and MSE as follows:

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \frac{1}{2N} \sum_{n=1}^N (y_n - \mathbf{x}_n^\top \mathbf{w})^2 \\ &= \frac{1}{2N} \mathbf{e}^\top \mathbf{e} \end{aligned}$$

then the gradient is given by

$$\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} \mathbf{X}^\top \mathbf{e}$$

Computational cost. What is the complexity (# operations) of computing the gradient?

a) starting from \mathbf{w} and

b) given \mathbf{e} and \mathbf{w} ?

$$\left. \begin{array}{l} \text{compute } \mathbf{e} \\ \overbrace{N \cdot D + N}^{\text{compute } \mathbf{e}} + \overbrace{N \cdot D + D}^{x_D^\top \cdot \frac{1}{N}} \end{array} \right\} O(N \cdot D)$$

Variant with offset. Recall: Alternative trick when also incorporating an offset term for the regression:

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad \tilde{\mathbf{X}} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1D} \\ 1 & x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \dots & x_{ND} \end{bmatrix}$$

$$f = w_0 + \tilde{x}_n^T \tilde{w}$$

$$= \tilde{x}_n^T \tilde{w}$$

$\nearrow \mathbb{R}^D$

$\nearrow \mathbb{R}^{D+1}$

Stochastic Gradient Descent

Sum Objectives. In machine learning, most cost functions are formulated as a **sum** over the training examples, that is

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w}) ,$$

where \mathcal{L}_n is the cost contributed by the n -th training example.

Q: What are the \mathcal{L}_n for linear MSE?

The SGD Algorithm. The stochastic gradient descent (SGD) algorithm is given by the following update rule, at step t :

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)}) .$$

$\in \mathbb{R}^D$

compare to GD!

Theoretical Motivation.

Idea: Cheap but unbiased estimate of the gradient!

In expectation over the random choice of n , we have

$$\mathbb{E} \left[\frac{\nabla \mathcal{L}_n(\mathbf{w})}{\mathbf{g}} \right] = \nabla \mathcal{L}(\mathbf{w})$$

which is the true gradient direction.
(check!)

$\mathbf{g} \in \mathbb{R}^D$
 \mathbf{g} is a stochastic gradient

Mini-batch SGD. There is an intermediate version, using the update direction being

$$\mathbf{g} := \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

iteration cost

$|B| \cdot D$

again with

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \mathbf{g} .$$

In the above gradient computation, we have randomly chosen a subset $B \subseteq [N]$ of the training examples. For each of these selected examples n , we compute the respective gradient $\nabla \mathcal{L}_n$, at the same current point $\mathbf{w}^{(t)}$.

The computation of \mathbf{g} can be parallelized easily. This is how current deep-learning applications utilize GPUs (by running over $|B|$ threads in parallel).

Note that in the extreme case $B := [N]$, we obtain (batch) gradient descent, i.e. $\mathbf{g} = \nabla \mathcal{L}$.

everyone among workers computes one of the $\nabla \mathcal{L}_n(w^{(t)})$ for $n \in B$

SGD for Linear MSE

See Exercise Sheet 2.

Computational cost. For linear MSE, what is the complexity (# operations) of computing the stochastic gradient?

(using only $|B| = 1$ data examples)

a) starting from \mathbf{w} ?

GD

$$\mathcal{L}(w) = \frac{1}{N} \sum_{n=1}^N \frac{1}{2} (y_n - x_n^\top w)^2$$

$$\underline{\nabla \mathcal{L}(w)} = -\frac{1}{N} X^\top (\underbrace{y - Xw}_{e})$$

SGD

$$\mathcal{L}_n(w) := \frac{1}{2} (y_n - x_n^\top w)^2$$

$$\underline{\nabla \mathcal{L}_n(w)} = \underbrace{(y_n - x_n^\top w)}_{\cdot} (-x_n)$$

$O(N \cdot D)$

$O(|B| \cdot D)$

$O(D)$

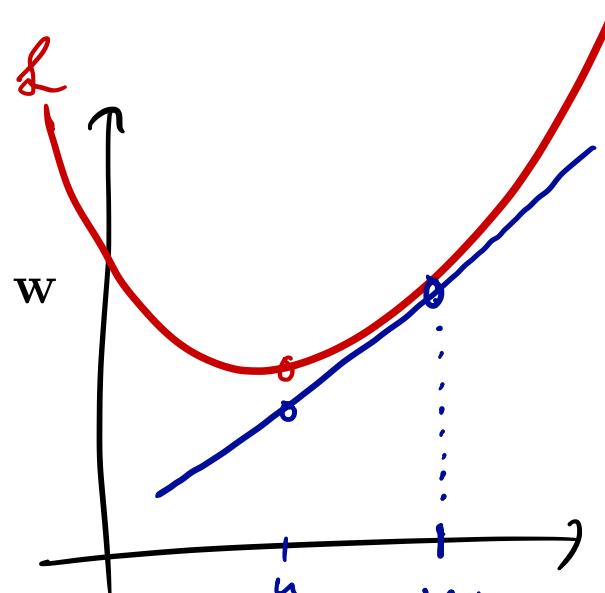
non-differentiable

Non-Smooth Optimization

An alternative characterization of *convexity*, for differentiable functions is given by

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \nabla \mathcal{L}(\mathbf{w})^\top (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}, \mathbf{w}$$

meaning that the function must always lie above its linearization.



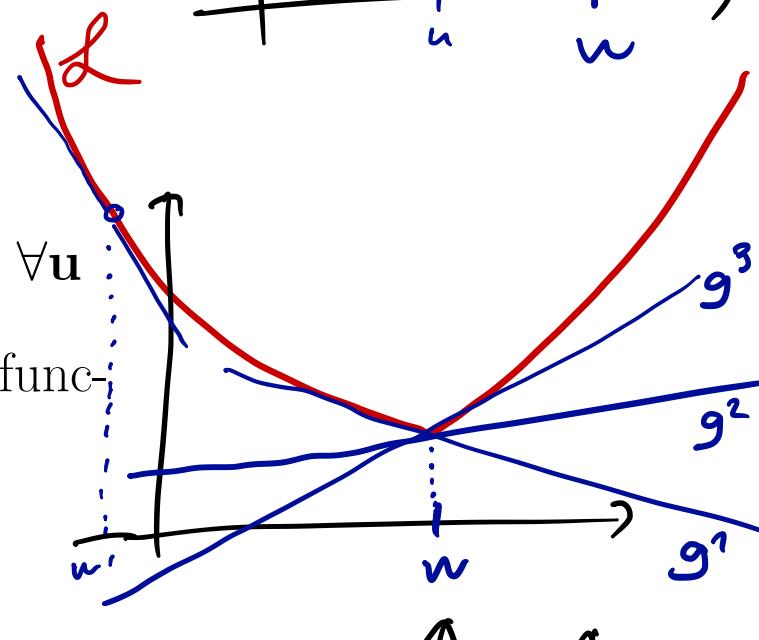
Subgradients

Def.

A vector $\mathbf{g} \in \mathbb{R}^D$ such that

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \mathbf{g}^\top (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}$$

is called a subgradient to the function \mathcal{L} at \mathbf{w} .



This definition makes sense for objectives \mathcal{L} which are not necessarily differentiable (and not even necessarily convex).

If \mathcal{L} is differentiable at \mathbf{w} , then the only subgradient at \mathbf{w} is $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$.

Subgradient Descent

Identical to the gradient descent algorithm, but using a subgradient instead of gradient. Update rule

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \mathbf{g}$$

for \mathbf{g} being a subgradient to \mathcal{L} at the current iterate $\mathbf{w}^{(t)}$.

Need:
 $\mathcal{L}(\mathbf{w})$ has a
 subgradient \mathbf{g}
 $\forall \mathbf{w} \in \mathbb{R}^D$

Example: Optimizing Linear MAE

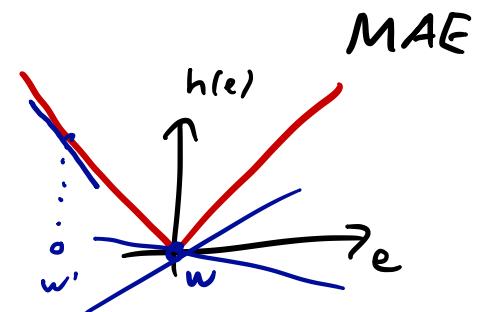
1. Compute a subgradient of the absolute value function

$$h : \mathbb{R} \rightarrow \mathbb{R}, h(e) := |e|.$$

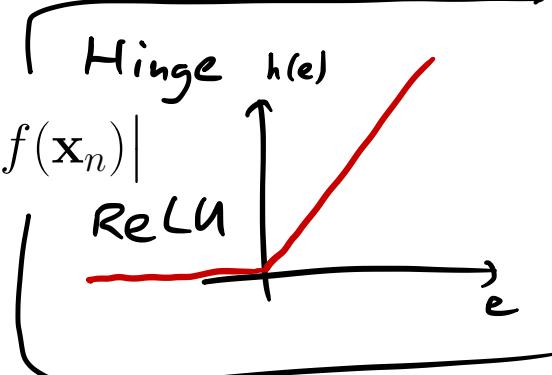
2. Recall the definition of the mean absolute error:

$$\mathcal{L}(\mathbf{w}) = \text{MAE}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$$

For linear regression, its (sub)gradient is easy to compute using the chain rule.
 Compute it!



$$g = \partial h(w) = \begin{cases} -1 & \text{if } e < 0 \\ [-1, 1] & \text{if } e = 0 \\ 1 & \text{if } e > 0 \end{cases}$$



use chain rule

See Exercise Sheet 2.

Stochastic Subgradient Descent

Stochastic SubGradient Descent
(still abbreviated SGD commonly).

Same, \mathbf{g} being a subgradient to the randomly selected \mathcal{L}_n at the current iterate $\mathbf{w}^{(t)}$.

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_n \mathcal{L}_n^{(w)}$$

Exercise: Compute the SGD update for linear MAE.

OVERVIEW

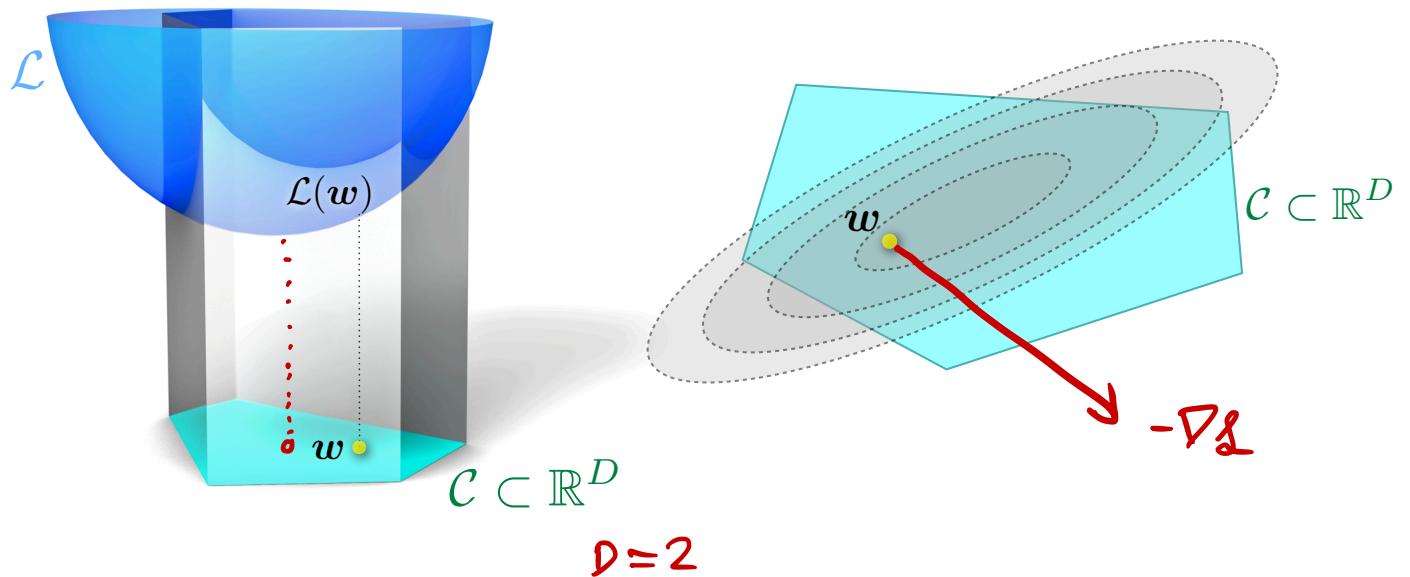
	Smooth	Non-Smooth
GD _g	$O(N \cdot D)$ MSE	$O(N \cdot D)$ MAE
SGD _{L_n}	$O(D)$ MSE	$O(D)$ MAE
(for linear regression)		

Constrained Optimization

Sometimes, optimization problems come posed with additional constraints:

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}), \quad \text{subject to } \mathbf{w} \in \mathcal{C}.$$

The set $\mathcal{C} \subset \mathbb{R}^D$ is called the constraint set.



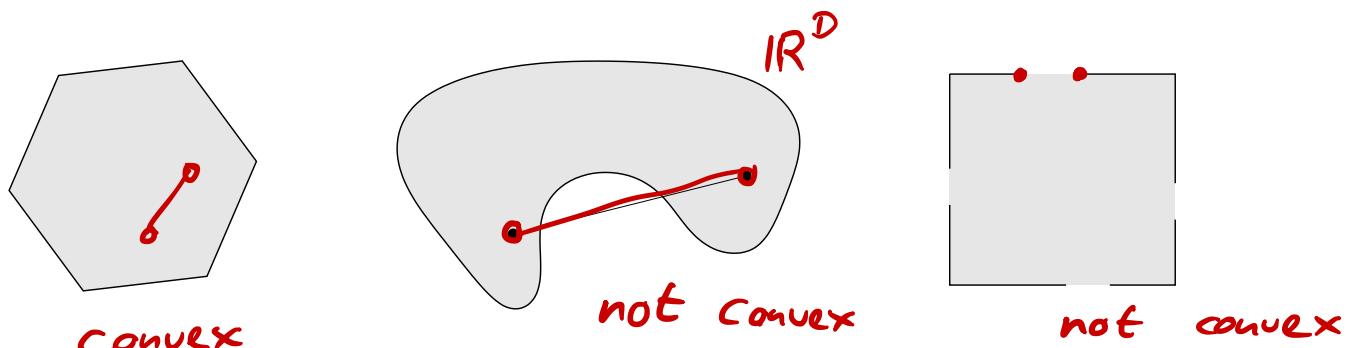
Solving Constrained Optimization Problems

- A) Projected Gradient Descent
- B) Transform it into an *unconstrained* problem

Convex Sets

A set \mathcal{C} is convex iff
the line segment between any two points of \mathcal{C} lies in \mathcal{C} , i.e., if for any $\mathbf{u}, \mathbf{v} \in \mathcal{C}$ and any θ with $0 \leq \theta \leq 1$, we have

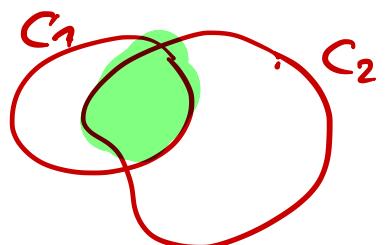
$$\theta\mathbf{u} + (1 - \theta)\mathbf{v} \in \mathcal{C}.$$



*Figure 2.2 from S. Boyd, L. Vandenberghe

Properties of Convex Sets.

- Intersections of convex sets are convex
- Projections onto convex sets are *unique*.
(and often efficient to compute)
Formal definition:
$$P_{\mathcal{C}}(\mathbf{w}') := \arg \min_{\mathbf{v} \in \mathcal{C}} \|\mathbf{v} - \mathbf{w}'\|.$$



A

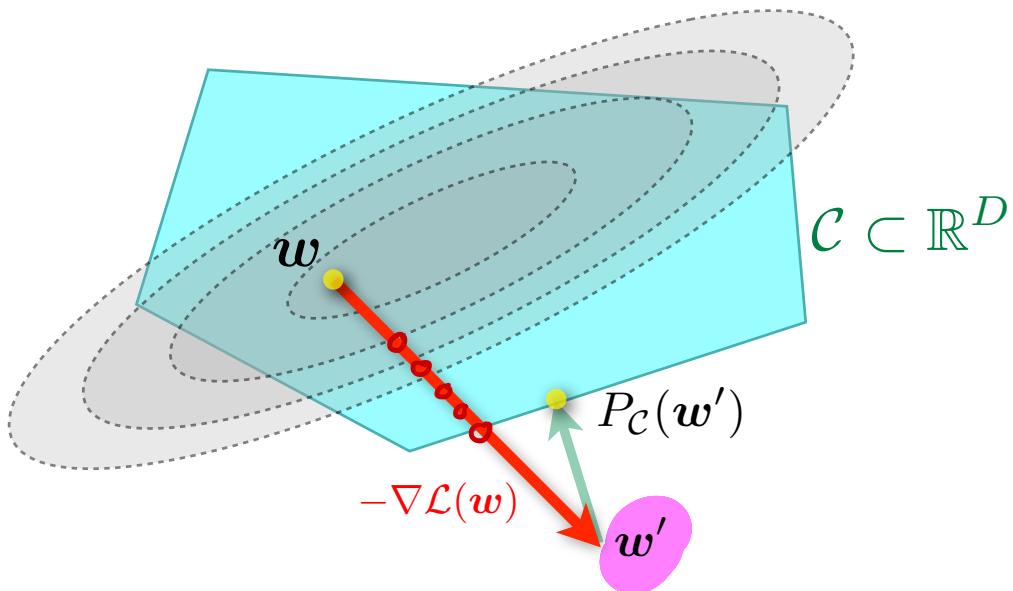
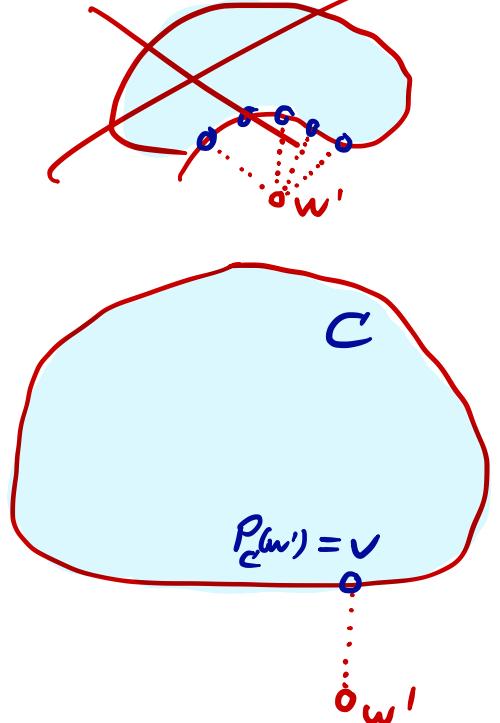
Projected Gradient Descent

Idea: add a projection onto \mathcal{C} after every step:

$$P_{\mathcal{C}}(\mathbf{w}') := \arg \min_{\mathbf{v} \in \mathcal{C}} \|\mathbf{v} - \mathbf{w}'\|.$$

Update rule:

$$\mathbf{w}^{(t+1)} := P_{\mathcal{C}}[\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})].$$



Projected SGD. Same SGD step, followed by the projection step, as above. Same convergence properties.

Computational cost of projection?
Crucial!

(B)

Turning Constrained into Unconstrained Problems

(Alternatives to projected gradient methods)

Use penalty functions instead of directly solving $\min_{\mathbf{w} \in \mathcal{C}} \mathcal{L}(\mathbf{w})$.

- “brick wall” (indicator function)

$$I_{\mathcal{C}}(\mathbf{w}) := \begin{cases} 0 & \mathbf{w} \in \mathcal{C} \\ \infty & \mathbf{w} \notin \mathcal{C} \end{cases}$$

$$\Rightarrow \min_{\mathbf{w} \in \mathbb{R}^D} \mathcal{L}(\mathbf{w}) + I_{\mathcal{C}}(\mathbf{w})$$

(disadvantage: non-continuous objective)

*m=1**A^{1x0}*

- Penalize error. Example:

$$\mathcal{C} = \{\mathbf{w} \in \mathbb{R}^D \mid \underbrace{A\mathbf{w} = \mathbf{b}}\}$$

$$\Rightarrow \min_{\mathbf{w} \in \mathbb{R}^D} \mathcal{L}(\mathbf{w}) + \lambda \underbrace{\|A\mathbf{w} - \mathbf{b}\|^2}_{\text{trade-off parameter}}$$

trade-off parameter

- Linearized Penalty Functions
(see Lagrange Multipliers)

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) + \mathbf{m}^T (\mathbf{A}\mathbf{w} - \mathbf{b})$$

$$\mathbf{m} \in \mathbb{R}^m$$

Implementation Issues

For Gradient Methods:

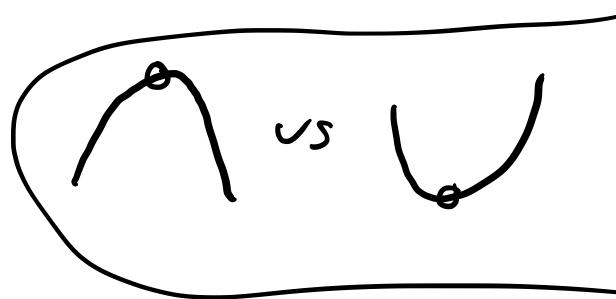
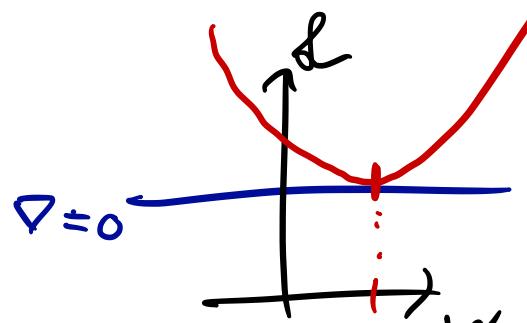
Stopping criteria: When $\nabla \mathcal{L}(\mathbf{w})$ is (close to) zero, we are (often) close to the optimum.

Optimality: If the second-order derivative is positive (positive semi-definite to be precise), then it is a (possibly local) minimum. If the function is also convex, then this condition implies that we are at a global optimum. See the supplementary section on **Optimality Conditions**.

Step-size selection: If γ is too big, the method might diverge. If it is too small, convergence is slow. Convergence to a local minimum is guaranteed only when $\gamma < \gamma_{min}$ where γ_{min} is a fixed constant that depends on the problem.

Unconstrained
opt

$$\nabla \mathcal{L}(\mathbf{w}) = 0$$



choose γ

(application
specific)

Line-search methods: For some objectives \mathcal{L} , we can set step-size automatically using a line-search method. More details on “backtracking” methods can be found in Chapter 1 of Bertsekas’ book on “nonlinear programming”.

Feature normalization: Gradient descent is very sensitive to ill-conditioning. Therefore, it is typically advised to normalize your input features. With unnormalized features, step-size selection becomes more difficult since different “directions” might converge at different “speed”.

see next
lecture!

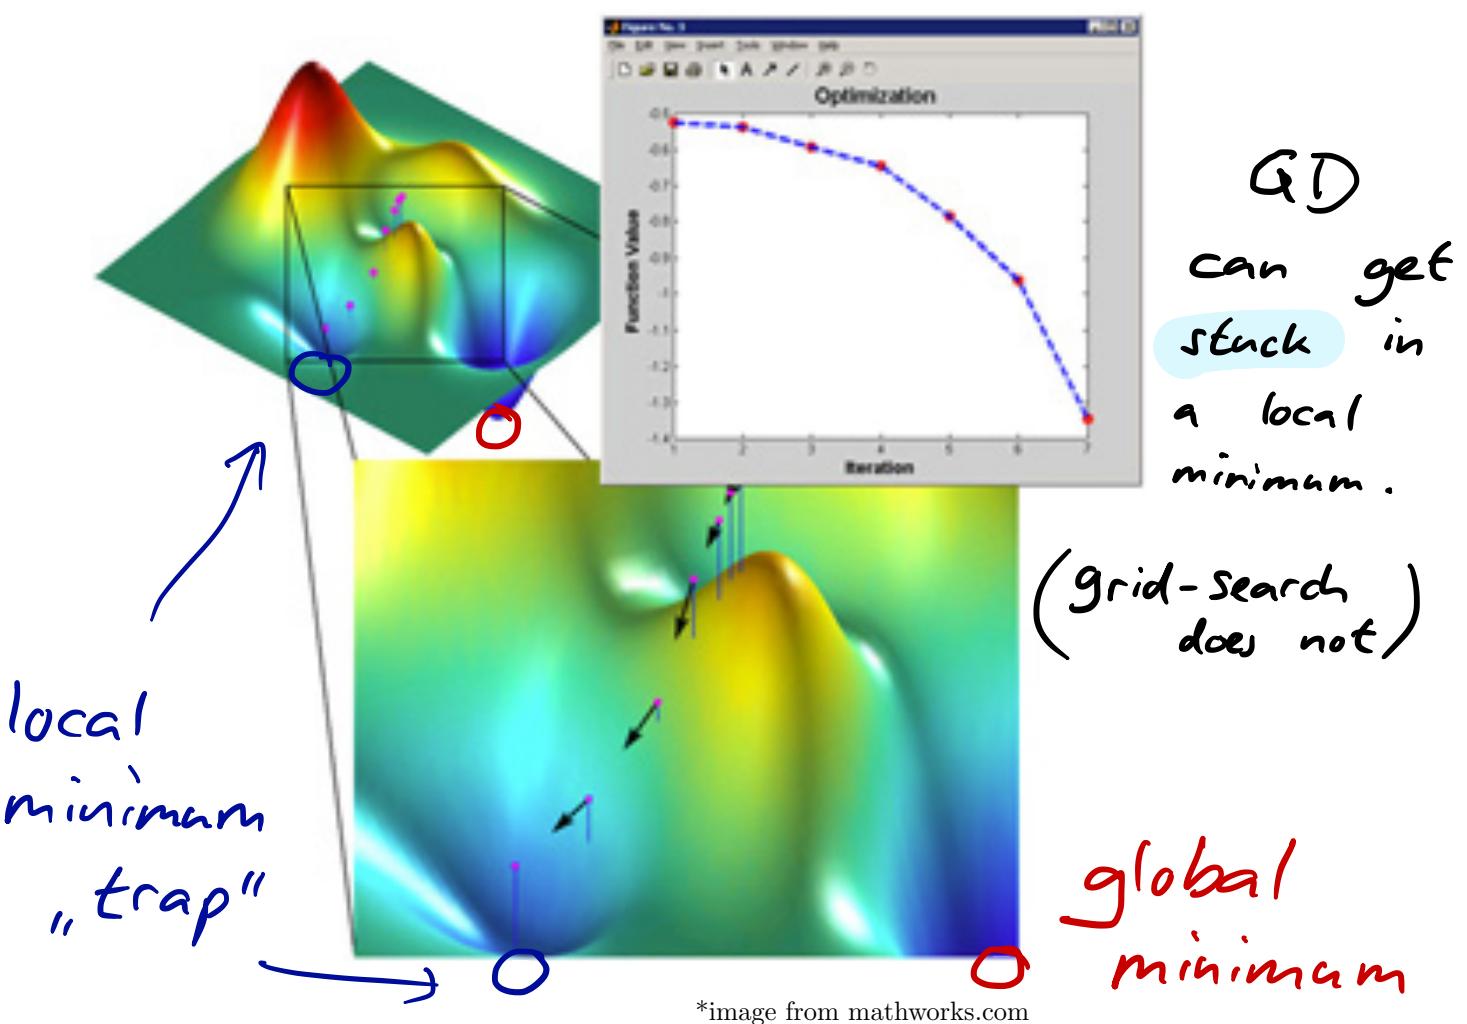
„fast“
some directions d with

$$\left| \frac{\partial}{\partial w_d} \mathcal{L} \right| \gg 0$$

„slow“
some directions d' with

$$\left| \frac{\partial}{\partial w_{d'}} \mathcal{L} \right| \approx 0$$

Non-Convex Optimization



Real-world problems are not convex!

All we have learnt on algorithm design and performance of convex algorithms still helps us in the non-convex world.

+ all examples we have seen so far are convex problems

Additional Notes

Grid Search and Hyper-Parameter Optimization

Read more about grid search and other methods for “hyperparameter” setting:

en.wikipedia.org/wiki/Hyperparameter_optimization#Grid_search.

Computational Complexity

The computation cost is expressed using the [big-O](#) notation. Here is a definition taken from Wikipedia. Let f and g be two functions defined on some subset of the real numbers. We write $f(x) = \mathcal{O}(g(x))$ as $x \rightarrow \infty$, if and only if there exists a positive real number c and a real number x_0 such that $|f(x)| \leq c|g(x)|$, $\forall x > x_0$.

Please read and learn more from this page in Wikipedia:

en.wikipedia.org/wiki/Computational_complexity_of_mathematical_operations#Matrix_algebra .

- What is the computational complexity of matrix multiplication?
- What is the computational complexity of matrix-vector multiplication?

Optimality Conditions

For a *convex* optimization problem, the first-order *necessary* condition says that at *an* optimum the gradient is equal to zero.

$$\nabla \mathcal{L}(\mathbf{w}^*) = \mathbf{0} \tag{1}$$

The second-order *sufficient* condition ensures that the optimum is a minimum (not a maximum or saddle-point) using the [Hessian](#) matrix,

which is the matrix of second derivatives:

$$\mathbf{H}(\mathbf{w}^*) := \frac{\partial^2 \mathcal{L}(\mathbf{w}^*)}{\partial \mathbf{w} \partial \mathbf{w}^\top} \quad \text{is positive semi-definite.} \quad (2)$$

The Hessian is also related to the convexity of a function: a twice-differentiable function is convex if and only if the Hessian is positive semi-definite at all points.

SGD Theory

As we have seen above, when N is large, choosing a random training example (\mathbf{x}_n, y_n) and taking an SGD step is advantageous:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \gamma^{(t)} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

For convergence, $\gamma^{(t)} \rightarrow 0$ “appropriately”. One such condition called the Robbins-Monroe condition suggests to take $\gamma^{(t)}$ such that:

$$\sum_{t=1}^{\infty} \gamma^{(t)} = \infty, \quad \sum_{t=1}^{\infty} (\gamma^{(t)})^2 < \infty \quad (3)$$

One way to obtain such sequence is $\gamma^{(t)} := 1/(1+t)^r$ where $r \in (0.5, 1)$.

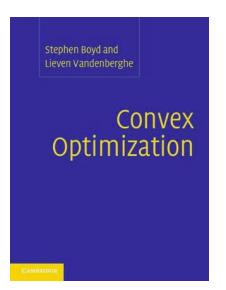
Read Section 8.5 of Kevin Murphy’s book. There are many variants of this method too. See more details and references in the Wikipedia page en.wikipedia.org/wiki/Stochastic_gradient_descent.

More Optimization Theory

If you want, you can gain a deeper understanding of several optimization methods relevant for machine learning from this survey:

[Convex Optimization: Algorithms and Complexity](#)
- by Sébastien Bubeck

And also from the [book of Boyd & Vandenberghe](#)
(both are free online PDFs)



(> 30 000
citations)

ToDo

1. Chain-rule



If it has been a while, familiarize yourself with it again.

2. Revise computational complexity (also see the Wikipedia link in Page 6 of lecture notes).
3. Derive the computational complexity of grid-search, gradient descent and stochastic gradient descent for linear MSE (# steps and cost per step).
4. Derive the gradients for the linear MSE and MAE cost functions.
5. Implement gradient descent and gain experience in setting the step-size.
6. Implement SGD and gain experience in setting the step-size.